

Ian Santos

San Mateo, CA | 650-515-2019 | ian6502011@gmail.com | github.com/TheBoiLeroy

PROFESSIONAL SUMMARY

Machine-Learning Engineer with hands-on expertise in **Large Language Model (LLM) fine-tuning, generative AI, prompt engineering, Retrieval-Augmented Generation (RAG), and MLOps**. Proven record translating state-of-the-art NLP research into production systems that improve user experience and broaden access to technology. Fluent in English and Spanish.

CORE SKILLS

LLMs & NLP	Transformers (GPT, BERT), Hugging Face, RAG, prompt engineering
ML Frameworks	PyTorch, TensorFlow/Keras, scikit-learn, spaCy
MLOps & DevOps	AWS, Docker, Kubernetes, CI/CD, GitHub Actions, MLflow, Weights&Biases, Flask-SocketIO
Programming	Python, TypeScript/Node.js, Dart/Flutter, React, C++, Java
Data & Databases	PostgreSQL, MySQL, MongoDB, SQLite; feature engineering; SQL/NoSQL

EXPERIENCE

Founding ML Engineer

Stealth Startup

Remote

Jan 2025 – Present

- Architected a **voice-driven investing assistant** converting speech to structured Alpaca trades via custom NER and intent classification.
- Built cross-platform app with Flutter and Supabase (postgres database)

Machine-Learning / Backend Engineer

Brinqa Sponsored Capstone, UCSC

Santa Cruz, CA

Sep 2024 – Dec 2024

- Led backend of a **Slack cybersecurity-risk chatbot**
- Developed TypeScript micro-services syncing Jira and Bitbucket APIs into an SQLite knowledge base
- Reviewed code contributions from engineering teammates for quality and consistency.
- Delivered MVP **one week early** within a 6-person Agile team.

Full-Stack Developer

Circle Jam Project, UCSC

Santa Cruz, CA

Jan 2024 – Mar 2024

- Launched React, Django, and PostgreSQL platform for listeners to compare stats and achievements.

Software Engineer

High-Concurrency HTTP Server, UCSC

Santa Cruz, CA

Sep 2023 – Dec 2023

- Designed C++ server handling via custom hash-locked cache and ACID-compliant logging.

AI Project Lead

Google Developer Student Club and AI Club

Santa Cruz, CA

Aug 2022

- Mentored peers on text generation, income prediction, and transfer learning projects using TensorFlow and Python.
- Conducted workshops on data preprocessing, model validation, and AI ethics.

SELECTED ACADEMIC PROJECTS

BERT Question-Answering Fine-Tune

CSE 142

UCSC

Winter 2024

- Implemented end-to-end BERT fine-tuning pipeline on SQuAD-v2 achieving **86 F1**.

Domain-Specific Sentiment Classifier

CSE 143

UCSC

Spring 2024

- Trained LSTM and fastText models, improving accuracy **15%** over baseline.

EDUCATION

University of California, Santa Cruz

B.S. Computer Science

Santa Cruz, CA

Dec 2024

Cañada College

A.S. Computer Science — Dean's List

Redwood City, CA

Jun 2020

ADDITIONAL INFORMATION

- **Languages:** English (native), Spanish (fluent)
- Volunteer ML mentor for first-generation students.